


Decode latency · host input → completed host output

NVIDIA GeForce RTX 4090 · float32 payloads · identical stdlib level-6 streams · source bb00aea62bae

Transfers, public API work, status checks and output conversion included; byte oracle and file I/O excluded.



Dots: all 12 samples per series. Bars and labels: median latency in ms. 2 warmups excluded; series interleaved in randomized rounds.

Memoryviews have different backing storage: CPU views decoded bytes; CUDA views a completed host NumPy array.

CPU: stdlib zlib. CUDA: synchronous host API. All outputs reside on the host. Each panel has its own latency scale.

Shared-workstation measurements; sample spread is displayed directly. These timings do not establish isolated kernel costs.