


Before / after codec optimisation

64 MiB inputs · warm resident API throughput · matching input and decode-stream hashes

● Before ◆ After

Compression

Each encoder's own output



Throughput (MiB/s) · log scale

Frozen-stream decoding

Identical baseline-encoded bytes



Throughput (MiB/s) · log scale

Stdlib level-6 decoding

Identical stdlib-encoded bytes



Throughput (MiB/s) · log scale

NVIDIA GeForce RTX 3090 · 30 samples per point combined from 2 runs per version (ABBA) · error bars: sample min-max, not confidence intervals

Warm synchronized API wall time includes allocations; startup, generation, transfers and correctness checks excluded.

CPU codec forbidden inside measured calls. CUDA-event kernel profiles are separate and are not plotted.