


Whole-process CUDA timeline


NVIDIA GeForce RTX 4090 · 64 MiB float32 · 20 validated measured iterations · source bb00aea62bae

Instrumented diagnostics · 2 warmups · synchronized panels share elapsed time; resource samples have lower temporal resolution

Host annotations · nested phases and repeated calls




Exact recorded CUDA intervals · stage colors match the Nsight figures



Copy bytes assigned at completion · stacked by transfer direction




Recorded kernel activity fraction · interval union / bin duration; not occupancy




GPU VRAM · process allocations and whole-device usage are separate scopes



Process-tree CPU · cumulative CPU-time differences at actual sample intervals



Process-tree resident memory · summed RSS can count shared pages more than once



Dashed lines bound the measured loop; whole-process shading locates it. CUDA bins: 10 ms; view-boundary bins use their actual clipped width.

Collector cadence: median 46.507 ms (44.824–56.325 ms). Metric points use query midpoints; CPU uses averaging-interval midpoints.

Host/CUDA alignment: $\pm 40.181 \mu\text{s}$ (half the origin-mark bracket). Host annotations do not establish GPU ownership or exclusive costs.

Transfers are completed bytes per bin, not instantaneous bandwidth. Coarse resource samples cannot resolve individual short phases; nulls remain gaps.

Nsight warns potentially missing NVTX events; all expected host phase ranges were validated. Other warnings remain in the report and manifest.